

SyriaTel Customer Churn Analysis and predictive Modeling

Introduction

For this analysis, we will use the dataset: [Churn in Telecom's dataset](#), which contains comprehensive information about churn, customer service calls, account length (the number of days a client has had their account active), total calls per day, daily charge, etc.

The goal is to help the telecom business reduce the amount of money lost due to customer churn and build a model that will help predict the type of client that will likely do so.

Methodology

Exploratory Data Analysis (EDA)

This step is essential to make sure things go smoothly and geting a good understanding of the dataset:

- dataset overview
- data cleaning
- data/ Business understanding

Statistical modeling

In this part we will use the insight that we get from the exploratory data analysis to make a predictive model of SyriaTel Customer churn.

Major questions

Questions that the management Team of SyriaTel might want to ask.These questions are based on our data and business understanding. They will be answered by our model. (recommendations will be provided along answering the major questions)

Summary

Contact information

✓ Exploratory Data Analysis (EDA)

✓ Dataset overview

```
1 import pandas as pd
2 df = pd.read_csv("/content/bigml_59c28831336c6604c800002a.csv")
3 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   state            3333 non-null    object  
 1   account length   3333 non-null    int64  
 2   area code         3333 non-null    int64  
 3   phone number     3333 non-null    object  
 4   international plan 3333 non-null    object  
 5   voice mail plan  3333 non-null    object  
 6   number vmail messages 3333 non-null    int64  
 7   total day minutes 3333 non-null    float64 
 8   total day calls   3333 non-null    int64  
 9   total day charge  3333 non-null    float64 
 10  total eve minutes 3333 non-null    float64 
 11  total eve calls   3333 non-null    int64  
 12  total eve charge  3333 non-null    float64 
 13  total night minutes 3333 non-null    float64 
 14  total night calls  3333 non-null    int64  
 15  total night charge 3333 non-null    float64 
 16  total intl minutes 3333 non-null    float64 
 17  total intl calls   3333 non-null    int64  
 18  total intl charge  3333 non-null    float64 
 19  customer service calls 3333 non-null    int64  
 20  churn             3333 non-null    bool  
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

✓ Data cleaning

```

1 # Create a copy of the dataset
2 df1 = df.copy()
3 # remove duplicate
4 df1.drop_duplicates(inplace=True)
5 # Remove spaces from column names for easier access.
6 df1.columns = df1.columns.str.replace(' ', '_')
7 # Set 'phone_number' as the index
8 df1.set_index('phone_number', inplace=True)
9 # we will drop area code, since there are about 51 states but just 3 area codes
10 df1.drop('area_code', axis=1, inplace=True)
11 # the dataset has no duplicates nor missing values
12 # the dataset was already pretty clean
13 # so many other cleaning steps won't be necessary.
14 # each remaining column also has their data in a proper data type.

```

▼ Data/ Business understanding

- check for class imbalance
- understanding

```

1 # let's check for class imbalance.
2 print('Raw counts: \n')
3 print(df1.churn.value_counts())
4 print('-----')
5 print('Normalized counts: \n')
6 normalized_counts = df['churn'].value_counts(normalize=True)
7 print (normalized_counts)
8 print(f"{normalized_counts[False]* 100:.0f}% of the result is false")

```

Raw counts:

```

churn
False    2850
True     483
Name: count, dtype: int64
-----
```

Normalized counts:

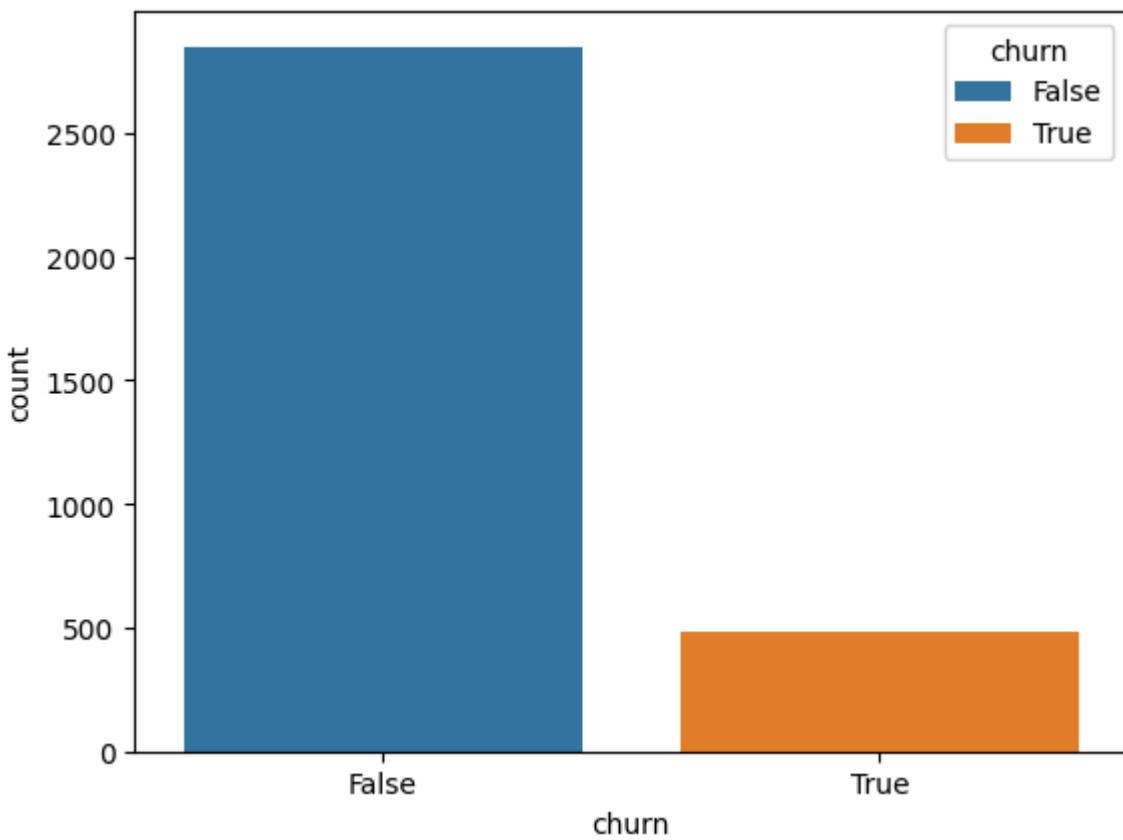
```

churn
False    0.855086
True     0.144914
Name: proportion, dtype: float64
86% of the result is false
```

```

1 # There is indeed class imbalance let's graph it.
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 sns.countplot(x='churn', data=df1, hue='churn')
5 # Save the plot to a file
6
7 plt.show()
8
```

```
9 # To save the plot as a JPEG file  
10 # plt.savefig('churn_countplot.jpeg')
```



Data Understanding

The dataset tracks SyriaTel customers' phone activity and contract details, including location, phone number, international and voicemail plans, and stored messages. It also records service usage—minutes, calls, and charges across day, evening, night, and international periods—plus customer service interactions. The target variable `churn` indicates whether a customer canceled their subscription. The goal is to identify patterns that predict churn.

Business Understanding

Customer churn poses a financial risk for SyriaTel, making retention a priority. Understanding why customers leave—whether due to service quality, pricing, or support—is key. By analyzing usage and subscription data, the company can identify at-risk customers and build a predictive model to guide proactive retention strategies.

❖ Statistical modeling

- baseline model
- perfecting model
- final model

each model will be followed by an analysis of its performance.

```
1 # to start we will import the necessary library
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler, OneHotEncoder
4 from sklearn.compose import ColumnTransformer
5 from sklearn.pipeline import Pipeline
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 import numpy as np
11 import warnings
12 from sklearn.exceptions import ConvergenceWarning
```

▼ Baseline model

we will use Logistic regression because it's simple, interpretable, and well-suited for binary classification.

```
1 # let's build a baseline model.
2 y = df1['churn']
3 X = df1.drop(['churn'], axis=1)
4 # Identify categorical and numerical columns
5 categorical_cols = X.select_dtypes(include=['object', 'category']).columns
6 numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns
7
8 # Apply one-hot encoding to categorical columns
9 X_categorical = pd.get_dummies(X[categorical_cols])
10
11 # Combine the numerical and encoded categorical columns
12 X = pd.concat([X[numerical_cols], X_categorical], axis=1)
13
14 # Split the data into training and test sets
15 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
1 # Suppress ConvergenceWarning from sklearn
2 warnings.filterwarnings("ignore", category=ConvergenceWarning, module='sklearn')
3
4 # Initialize the logistic regression Classifier with imbalanced class weights
5 logreg = LogisticRegression(fit_intercept=False, solver='lbfgs')
6
7 # Fit the model
8 model_log = logreg.fit(X_train, y_train)
9 print(model_log)
10
11 # Predict probabilities on the test set
12 y_score = logreg.decision_function(X_test)
13
14 # Calculate ROC curve and AUC (needed for AUC printout)
15 fpr, tpr, thresholds = roc_curve(y_test, y_score)
```

```

16 roc_auc = auc(fpr, tpr)
17
18 print(f'AUC for imBalanced Class Weight: {roc_auc}')
19 print('-----')
20
21 # Make predictions for classification report and confusion matrix
22 y_pred = logreg.predict(X_test)
23
24 # Print Classification Report
25 print("\nClassification Report (imBalanced Class Weight):")
26 print(classification_report(y_test, y_pred))
27
28 # Print Confusion Matrix
29 print("\nConfusion Matrix (imBalanced Class Weight):")
30 print(confusion_matrix(y_test, y_pred))

```

```

LogisticRegression(fit_intercept=False)
AUC for imBalanced Class Weight: 0.7112656467315717
-----
```

	precision	recall	f1-score	support
False	0.87	0.97	0.92	719
True	0.33	0.09	0.14	115
accuracy			0.85	834
macro avg	0.60	0.53	0.53	834
weighted avg	0.80	0.85	0.81	834

```

Confusion Matrix (imBalanced Class Weight):
[[699  20]
 [105  10]]
```

Analysis

the baseline model performs well enough to predict non churn but very poorly on churn (with low precision, accuracy and F-1 score) . and since we are mostly interested in churn we will keep tuning the model.

perfecting model.

we will have the baseline model tuned differently so that we can find the best model possible. In the end all model will be represented on a ROC curve for a better comparaison.

Model 2

```

1 # Suppress ConvergenceWarning from sklearn
2 warnings.filterwarnings("ignore", category=ConvergenceWarning, module='skle
3
4 # Initialize the logistic regression
5 logreg = LogisticRegression(fit_intercept=False, solver='lbfgs', class_weig
6
7 # Fit the model
8 model_log = logreg.fit(X_train, y_train)
9 # Predict probabilities on the test set
10 y_score = logreg.decision_function(X_test)
11 y_predict = logreg.predict(X_test) # Predictions for Model 2
12
13 # Calculate ROC curve and AUC for Model 2
14 fpr_model2, tpr_model2, thresholds_model2 = roc_curve(y_test, y_score)
15 roc_auc_model2 = auc(fpr_model2, tpr_model2)
16
17 print(f'AUC for Model 2: {roc_auc_model2}')
18 print('-----')
19
20 # Print Classification Report
21 print("\nClassification Report (Model 2):")
22 print(classification_report(y_test, y_predict))
23
24 # Print Confusion Matrix
25 print("\nConfusion Matrix (Model 2):")
26 print(confusion_matrix(y_test, y_predict))

```

AUC for Model 2: 0.7385862006409869

	precision	recall	f1-score	support
False	0.92	0.70	0.80	719
True	0.26	0.64	0.37	115
accuracy			0.69	834
macro avg	0.59	0.67	0.58	834
weighted avg	0.83	0.69	0.74	834

Confusion Matrix (Model 2):
[[503 216]
 [41 74]]

Analysis

This model showed improved performance in identifying churn, with recall for True rising from 0.09 in the baseline to 0.64 in the model. It also has a higher F1-score for the True class (the customer actually churn).

Model 3

```

1 # Suppress ConvergenceWarning from sklearn
2 warnings.filterwarnings("ignore", category=ConvergenceWarning, module='skle
3
4 # Initialize the logistic regression
5 logreg = LogisticRegression(fit_intercept=False, solver='lbfgs', class_wei
6
7 # Fit the model
8 model_log = logreg.fit(X_train, y_train)
9 # Predict probabilities on the test set
10 y_score = logreg.decision_function(X_test)
11 y_pred = logreg.predict(X_test)
12
13 # Calculate ROC curve and AUC for Model 3
14 fpr_model3, tpr_model3, thresholds_model3 = roc_curve(y_test, y_score)
15 roc_auc_model3 = auc(fpr_model3, tpr_model3)
16
17 print(f'AUC for Model 3: {roc_auc_model3}')
18 print('-----')
19 # Print Classification Report
20 print("\nClassification Report (Model 3):")
21 print(classification_report(y_test, y_pred))
22
23 # Print Confusion Matrix
24 print("\nConfusion Matrix (Model 3):")
25 cm3 = confusion_matrix(y_test, y_pred)
26 print(cm3)

```

AUC for Model 3: 0.8067243151720385

Classification Report (Model 3):

	precision	recall	f1-score	support
False	0.95	0.74	0.83	719
True	0.31	0.75	0.44	115
accuracy			0.74	834
macro avg	0.63	0.74	0.64	834
weighted avg	0.86	0.74	0.78	834

Confusion Matrix (Model 3):

```

[[531 188]
 [ 29  86]]

```

Analysis

this model performs better in precision for both True and False and the f1-score has come from 0.37 to 0.44. this model is better than the two former.

Model 4

```

1 # Import the XGBoost library for gradient boosting
2 import xgboost as xgb
3 # Initialize the XGBoost Classifier
4 # Use scale_pos_weight to handle class imbalance. This parameter is used to
5 # in the training data by giving more weight to the minority class (churned)
6 churn_counts = df1['churn'].value_counts()
7 # Calculate the ratio of non-churned to churned customers
8 scale_pos_weight = churn_counts[False] / churn_counts[True]
9 xgb_model = xgb.XGBClassifier(objective='binary:logistic', # Set the object
10                                eval_metric='logloss', # Set the evaluation metric
11                                use_label_encoder=False, # Suppress the warning
12                                scale_pos_weight=scale_pos_weight, # Apply the weight
13                                random_state=0) # Set the random state for reproducibility
14
15 # Fit the model
16 xgb_model.fit(X_train, y_train)
17
18 # Predict probabilities on the test set
19 # We are interested in the probability of the positive class (churn=True),
20 y_score_xgb = xgb_model.predict_proba(X_test)[:, 1]
21
22 # Calculate ROC curve and AUC
23 fpr_xgb, tpr_xgb, thresholds_xgb = roc_curve(y_test, y_score_xgb)
24 roc_auc_xgb = auc(fpr_xgb, tpr_xgb)
25
26 # Print the AUC score
27 print(f'AUC for XGBoost Model: {roc_auc_xgb}')
28 print('-----')
29
30 # Make predictions for classification report and confusion matrix
31 y_pred_xgb = xgb_model.predict(X_test)
32 # Print Classification Report
33 print("\nClassification Report (XGBoost Model):")
34 print(classification_report(y_test, y_pred_xgb))
35

```

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [0]

Parameters: { "use_label_encoder" } are not used.

```

bst.update(dtrain, iteration=i, fobj=obj)
AUC for XGBoost Model: 0.9315353449839754
-----
```

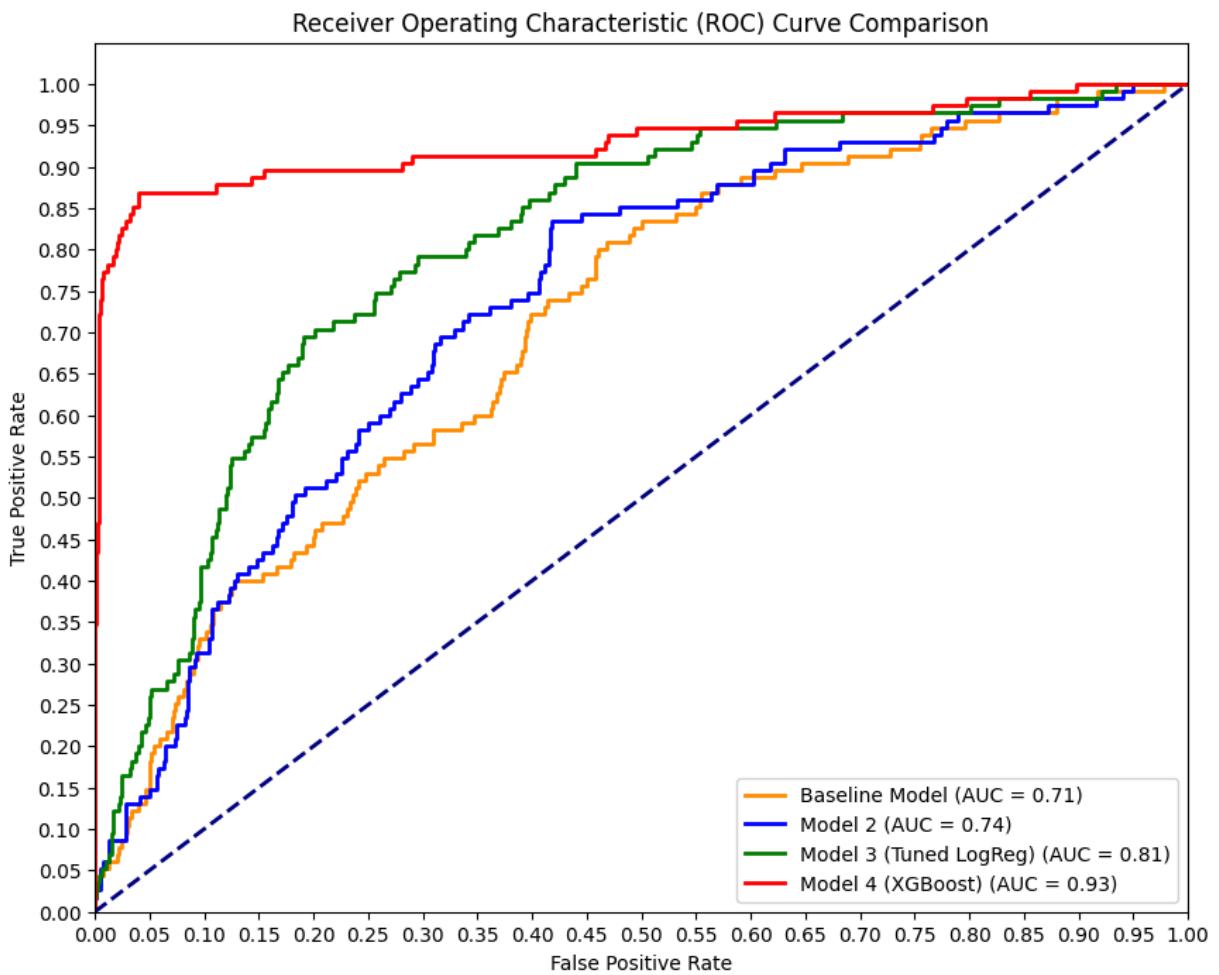
Classification Report (XGBoost Model):				
	precision	recall	f1-score	support
False	0.97	0.97	0.97	719
True	0.81	0.83	0.82	115
accuracy			0.95	834
macro avg	0.89	0.90	0.90	834
weighted avg	0.95	0.95	0.95	834



Analysis

this model perform great overall whether it's precision or recall. it is also likely to perform well on unseen data. with an AUC of 0.93 it's the best of all the models which is why we will keep it.

```
1 # Plot ROC curves for all models on the same graph
2 plt.figure(figsize=(10, 8))
3 lw = 2
4
5 # Plot ROC for Baseline Model (Model 1)
6 # fpr and tpr for the baseline model
7 plt.plot(fpr, tpr, color='darkorange', lw=lw, label=f'Baseline Model (AUC =
8 # Model 2
9 plt.plot(fpr_model2, tpr_model2, color='blue', lw=lw, label=f'Model 2 (AUC
10
11
12 # Plot ROC for Model 3
13 plt.plot(fpr_model3, tpr_model3, color='green', lw=lw, label=f'Model 3 (Tur
14
15 # Plot ROC for Model 4 (XGBoost Model)
16 plt.plot(fpr_xgb, tpr_xgb, color='red', lw=lw, label=f'Model 4 (XGBoost) (A
17
18
19 # Plot the diagonal line
20 plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
21
22 # Set plot limits and labels
23 plt.xlim([0.0, 1.0])
24 plt.ylim([0.0, 1.05])
25 plt.yticks([i/20.0 for i in range(21)])
26 plt.xticks([i/20.0 for i in range(21)])
27 plt.xlabel('False Positive Rate')
28 plt.ylabel('True Positive Rate')
29 plt.title('Receiver Operating Characteristic (ROC) Curve Comparison')
30 plt.legend(loc='lower right')
31 plt.show()
```



▼ Major questions

bellow are questions that the management Team of SyriaTel might want to ask.These questions are based on our data and business understanding .They will be answered by our model.

- What customers are most likely to churn ?
- Do customer service calls play a significant role in churn?
- Based on our analysis, what are actionable strategies to reduce churn?
- Which specific states have the highest churn rates?

▼ What customers are most likely to churn ?

```

1 # Analyze feature importance from the XGBoost model
2 # The feature_importances_ attribute provides the importance score for each
3 feature_importances = xgb_model.feature_importances_
4
5 # Get the names of the features
6 features = X_train.columns
7
8 # Create a pandas Series for easier handling and sorting

```

```

9 feature_importance_series = pd.Series(feature_importances, index=features)
10
11 # Sort the features by importance in descending order
12 sorted_feature_importance = feature_importance_series.sort_values(ascending=
13
14 # Print the top 10 most important features
15 print("Top 10 Most Important Features for Churn Prediction (XGBoost):")
16 print(sorted_feature_importance.head(10))

```

```

Top 10 Most Important Features for Churn Prediction (XGBoost):
international_plan_no      0.100623
customer_service_calls     0.071151
state_SD                   0.059375
state_MN                   0.051422
state_VA                   0.044444
state_FL                   0.040789
total_day_minutes          0.033859
state_ID                   0.033263
number_vmail_messages      0.031723
state_ME                   0.031024
dtype: float32

```

Note: the customers that make the most international plan along with the most customer service calls are the most likely to churn.

Recommendation: the company may want to do additional data gathering on international calls, and try to see which of it's aspect botters client.

▼ Do customer service calls play a significant role in churn?

Customers that has the most customer service call are the most likely to churn. This is a sign that the customer service might still has room for improvement.

Recommendation: the company should put more effort on understanding clients problem, and fix them as effectively as possible. the call from customers should be handle with tact, as customer satisfaction is likely to reduce churn.

Based on our analysis, what are actionable strategies to reduce churn?

to reduce churn the company might want to give the clients a better experience on international calls. they might also want to handle customer call more efficiently. since clients with most international and customer service call churn.

Which specific states have the highest churn rates? Mariland, Minesota and Virginia are the states that have the highest churn Rates.

Recommendation: further data should be gathered on this specific states. Poles might be a great way to proceed.

Summary

For this analysis, we used the [Churn in Telecom's dataset](#). We began with a dataset overview and performed necessary data cleaning. After cleaning, we gained a good understanding of the data and the business problem it addresses. We then developed different models to find the one that provides the best predictions on unseen data. Finally, we addressed major questions based on our analysis and provided recommendations.

▼ Contact Information

- First Name: Haender Michael
- Last Name: Jean Louis